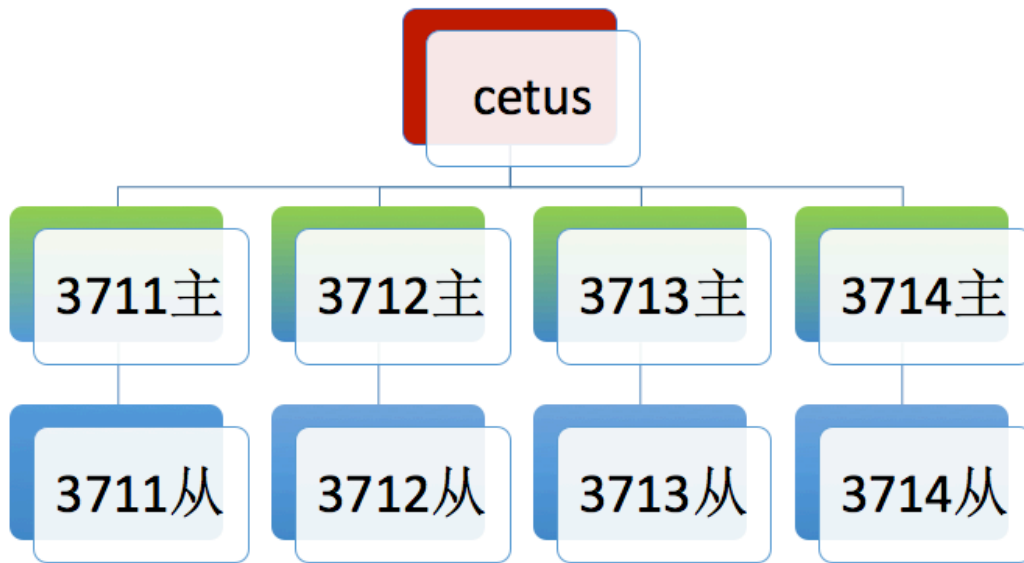


MySQL 5.7.21 XA bug 的详细验证

发现的 XA bug 现象是数据库的 xa 状态跟 binlog 记录的状态不一致。测试 MySQL 版本为 5.7.21。连接分片中间件，运行模拟银行转账的脚本，测试过程随机 kill -9 杀掉后端数据库，过 10 分钟后，查看各个后端数据库的 xa 悬挂事务，并且比对分片 cetus 的 xa log 和 MySQL binlog，判断悬挂事务的是否正常。一个 cetus 后面接四个写节点 3711/3712/3713/3714，每个写节点带一个从库：



分两种情况分析已发现的 xa bug

主库存在 xa 事务：

3713 主库发现 xa 事务，其他主库和所有从库未发现与之对应的 xa 事务

```
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 21
Server version: 5.7.21-log MySQL Community Server (GPL)

Copyright (c) 2000, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> xa recover;
+-----+-----+-----+-----+
| formatID | gtrid_length | bqual_length | data |
+-----+-----+-----+-----+
| 1 | 32 | 0 | clt-0.0.0.0:7001-5496_23_1229736 |
| 1 | 32 | 0 | clt-0.0.0.0:7001-5496_23_1230780 |
+-----+-----+-----+-----+
2 rows in set (0.00 sec)

mysql>
```

检查 clt-0.0.0.0:7001-5496_23_1229736 在 cetus 中的执行情况，可以看到 cetus 已经向四个写节点发送 xa commit，这个事务在 cetus 看来，已经成功了。

```
[mysql-cetus@hz-10- [REDACTED] logs6001]$ cat xa.log_23|grep clt-0.0.0.0:7001-5496_23_1229736
2018/04/10 23:30:03 +508 [info] XA COMMIT 'clt-0.0.0.0:7001-5496_23_1229736' 10. [REDACTED]:3711@43,10. [REDACTED]:3712@25,10. [REDACTED]:3713@65,10. [REDACTED]:3714@3
```

提示：xa recover 看到的 data，在 binlog 中记录的是十六进制的小写形式，可以在 python 中做一下转换

python

```
import base64
```

```
print base64.b16encode(b'clt-0.0.0.0:7001-5496_23_1229736').lower()
```

```
[mysql-cetus@hz-10- [REDACTED] logs6001]$ python
import base64
Python 2.7.5 (default, Oct 11 2015, 17:47:16)
[GCC 4.8.3 20140911 (Red Hat 4.8.3-9)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import base64
>>> print base64.b16encode(b'clt-0.0.0.0:7001-5496_23_1229736').lower()
636c742d302e302e302e303a373030312d353439365f32335f31323239373336
```

再看四个主库和四个从库的 binlog

3711 主库：

主库 binlog 已经记录 xa commit，并且实例中没有悬挂事务

```
[mysql3711@ [REDACTED] data]$ mysqlbinlog -vvv mysql-bin.000012> a.txt
[mysql3711@ [REDACTED] data]$ cat a.txt |grep 636c742d302e302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013711 end_log_pos 1441946 CRC32 0x0858e931 XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

3711 从库：

现象同 3711 主库，已 xa commit，无悬挂事务

```
[mysql3711@ [REDACTED] data]$ mysqlbinlog -vvv mysql-bin.000003> a.txt
[mysql3711@ [REDACTED] data]$ cat a.txt |grep 636c742d302e302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013711 end_log_pos 11838132 CRC32 0x01ae9fb1 XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

3712 主库：

现象同 3711 主库，已 xa commit，无悬挂事务

```
[mysql3712@centos7 ~]$ cat a.txt |grep 636c742d302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013712 end_log_pos 1464924 CRC32 0x857c2ae1 XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

3712 从库：

现象同 3711 主库，已 xa commit，无悬挂事务

```
[mysql3712@centos7 ~]$ cat a.txt |grep 636c742d302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013712 end_log_pos 11880723 CRC32 0xe984df44 XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

3713 主库：

主库 binlog 已经记录 xa commit，但是出现了悬挂事务！下面在看看 3713 从库

```
[mysql3713@centos7 ~]$ cat a.txt |grep 636c742d302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013713 end_log_pos 11946257 CRC32 0xb85804c6 XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

3713 从库：

现象同 3711 主库，已 xa commit，无悬挂事务，证明 3713 主库的 binlog 已同步到 3713 从库，并且从库 slave_sql_thread 已经成功应用传过来的 binlog

```
[mysql3713@centos7 ~]$ cat a.txt |grep 636c742d302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013713 end_log_pos 11946257 CRC32 0xb85804c6 XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

3714 主库：

现象同 3711 主库，已 xa commit，无悬挂事务

```
[mysql3714@centos7 ~]$ cat a.txt |grep 636c742d302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013714 end_log_pos 1446776 CRC32 0x2438eab0 XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

3714 从库：

现象同 3711 主库，已 xa commit，无悬挂事务

```
[mysql3714@centos7 ~]$ cat a.txt |grep 636c742d302e302e303a373030312d353439365f32335f31323239373336
XA START X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA END X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
#180410 23:30:03 server id 1013714 end_log_pos 1446776 CRC32 0x2438eab0 XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA PREPARE X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
XA COMMIT X'636c742d302e302e303a373030312d353439365f32335f31323239373336',X'',1
```

结论：cetus 向所有后端写节点发送 xa commit，3713 主库接收到 xa commit 命令，将 xa commit 命令写入本实例 binlog，binlog 被拉到从库，主库执行 xa commit 前被用 kill-9 杀掉。主库重新启动后不能识别 xa 事务的正确状态，导致了悬挂事务。

从库存在 xa 事务：

3714 从库发现 xa 事务，其他主库和所有从库未发现与之对应的 xa 事务

```

Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 243
Server version: 5.7.21-log MySQL Community Server (GPL)

Copyright (c) 2000, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> xa recover;
+-----+-----+-----+-----+
| formatID | gtrid_length | bqual_length | data |
+-----+-----+-----+-----+
| 1 | 32 | 0 | clt-0.0.0.0:7001-5496_23_1227527 |
+-----+-----+-----+-----+
1 row in set (0.00 sec)

```

检查 clt-0.0.0.0:7001-5496_23_1227527 在 cetus 中的执行情况，可以看到 cetus 发到 3714 主库的 xa prepare 失败，然后回滚另外三个写节点的 xa prepare。

```

2018/04/10 23:29:04 +145 [warn] XA PREPARE 'clt-0.0.0.0:7001-5496_23_1227527' 10.120.12.84:3714@9 failed
2018/04/10 23:29:04 +145 [info] XA ROLLBACK 'clt-0.0.0.0:7001-5496_23_1227527' 10.120.12.84:3711@22,10.120.12.84:3712@4,10.120.12.84:3713@62

```

转换 clt-0.0.0.0:7001-5496_23_1227527 为十六进制

```

import base64
Python 2.7.5 (default, Oct 11 2015, 17:47:16)
[GCC 4.8.3 20140911 (Red Hat 4.8.3-9)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import base64
>>> print base64.b16encode(b'clt-0.0.0.0:7001-5496_23_1227527').lower()
636c742d302e302e302e303a373030312d353439365f32335f31323237353237

```

3711 主库:

主库 binlog 已经记录 xa prepare 和 xa rollback，并且实例中没有悬挂事务

```

XA START X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server id 1013711 end_log_pos 195894 CRC32 0x423b5a00 XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA ROLLBACK X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1

```

3711 从库:

同 3711 主库，已经记录 xa prepare 和 xa rollback，并且实例中没有悬挂事务

```

XA START X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server id 1013711 end_log_pos 9170603 CRC32 0x0263cccc XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA ROLLBACK X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1

```

3712 主库:

同 3711 主库，已经记录 xa prepare 和 xa rollback，并且实例中没有悬挂事务

```

XA START X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server id 1013712 end_log_pos 229610 CRC32 0x8b0d1f06d XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA ROLLBACK X'636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1

```

3712 从库:

同 3711 从库，已经记录 xa prepare 和 xa rollback，并且实例中没有悬挂事务

```
XA START '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server_id 1013712 end_log_pos 9184720 CRC32 0xb7fd4eb4 XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA ROLLBACK '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
```

3713 主库:

同 3711 主库，已经记录 xa prepare 和 xa rollback，并且实例中没有悬挂事务

```
XA START '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server_id 1013713 end_log_pos 9497483 CRC32 0xd10204e0 XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA ROLLBACK '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
```

3713 从库:

同 3711 从库，已经记录 xa prepare 和 xa rollback，并且实例中没有悬挂事务

```
XA START '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server_id 1013713 end_log_pos 9275612 CRC32 0xdd377bae XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA ROLLBACK '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
```

3714 主库:

主库只记录了 xa prepare，没有 xa rollback，但是实例中没有悬挂事务

```
XA START '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server_id 1013714 end_log_pos 253269 CRC32 0xa1ca706c XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
```

3714 从库:

从库同样只记录了 xa prepare，没有 xa rollback，但是实例中有悬挂事务

```
XA START '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA END '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
#180410 23:29:04 server_id 1013714 end_log_pos 9182496 CRC32 0x2bd24386 XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
XA PREPARE '636c742d302e302e302e303a373030312d353439365f32335f31323237353237',X'',1
```

结论：cetus 向写节点发送 xa prepare，mysql3714 接收到 xa prepare 命令，将 xa prepare 命令写入本实例 binlog，binlog 被拉到从库，主库执行 xa prepare 前被 kill -9 杀掉，主库重新启动不能正确识别 xa 事务的正确状态，导致主库没有悬挂事务，而从库有悬挂事务，这种情况需要把从库 xa 事务进行 rollback。